

PART B - Classification

```
In [2]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import utils
import random
import os
import collections
import numpy
from sklearn import svm, tree
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import TruncatedSVD
from IPython.core.display import display, HTML
```

```
In [3]: # Load dataset from saved file
DATASET_NAME = 'broad'
dataset = utils.load_from_json('%s.json' % DATASET_NAME)
N = len(dataset[dataset.keys()[0]]) # Number of sounds per class
CLASS_NAMES = dataset.keys()
print 'Loaded dataet "%s" (%i classes, %i sounds per class)' % (DATASET_
NAME, len(dataset.keys()), N)
```

Loaded dataet "broad" (3 classes, 100 sounds per class)

```
In [4]: # Set configuration parameters (to be used below)
NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR = 10 # Maximum number of dimensio
ns for the feature vector.
CLASSIFIER_TYPE = 'tree' # Use 'svm' or 'tree'
PERCENTAGE_OF_TRAINING_DATA = 0.5 # Percentage of sounds that will be us
ed for training (others are for testing)
MAX_INPUT_TAGS_FOR_TESTING = 50 # Use a big number to "omit" this parame
ter and use as many tags as originally are in the sound
DESCRIPTORS = False
NO_KEYWORDS = ['http', 'rel="nofollow"', 'href', 'https', 'sound' , 'rec
orded', 'This', 'sounds']
```

3) Define vector space

```

In [5]: def build_tag_vector_space(n_dimensions, dataset, class_names):
        # Get all tags in the dataset (the vocabulary)
        all_tags = list()
        for class_name in class_names:
            class_tags = utils.get_all_tags_from_class(class_name, dataset)
            all_tags += class_tags

        all_tags = [tag for tag in all_tags if tag not in NO_KEYWORDS]
        # Filter out tags with less frequency (get only top N tags)
        most_common_tags = [tag for tag, count in collections.Counter(all_tags).most_common(n_dimensions)]
        filtered_tags = [tag for tag in most_common_tags if tag in all_tags]

        # Build our prototype feature vector (unique list of tags), and print first 10 tags
        prototype_feature_vector = list(set(filtered_tags))
        print 'Created prototype feature vector with %i dimensions (original %i dimensions)' % (
            len(prototype_feature_vector), len(set(all_tags)))
        print 'Prototype vector tags (sorted by occurrence in filtered_tags):', ', '.join([tag for tag in filtered_tags[:10]]),
        print '...\n' if len(filtered_tags) > 10 else '\n'
        return prototype_feature_vector

    def build_keyword_vector_space(n_dimensions, dataset, class_names):
        # Get all keywords in the dataset (the vocabulary)
        all_keywords = list()
        for class_name in class_names:
            class_keywords = utils.get_all_keywords_from_class(class_name, dataset)
            all_keywords += class_keywords

        all_keywords = [k for k in all_keywords if k not in NO_KEYWORDS]
        # Filter out keywords with less frequency (get only top N keywords)
        most_common_keywords = [k for k, count in collections.Counter(all_keywords).most_common(n_dimensions)]
        filtered_keywords = [k for k in most_common_keywords if k in all_keywords]

        # Build our prototype feature vector (unique list of keywords), and print first 10 keywords
        prototype_feature_vector = list(set(filtered_keywords))
        print 'Created prototype feature vector with %i dimensions (original %i dimensions)' % (
            len(prototype_feature_vector), len(set(all_keywords)))
        print 'Prototype vector keywords (sorted by occurrence in filtered_keywords):', ', '.join([k for k in filtered_keywords[:10]]),
        print '...\n' if len(filtered_keywords) > 10 else '\n'
        return prototype_feature_vector

    prototype_feature_vector_tags = build_tag_vector_space(
        n_dimensions=NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR,
        dataset=dataset,
        class_names=CLASS_NAMES,
    )

    prototype_feature_vector_keywords = (build_keyword_vector_space(
        n_dimensions=NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR,
        dataset=dataset,
        class_names=CLASS_NAMES,
    ))

    prototype_feature_vector = list(set(prototype_feature_vector_tags) | set(prototype_feature_vector_keywords) - set(NO_KEYWORDS))
    print 'Final feature vector size ' + str(len(prototype_feature_vector)) + ' features: ' + str(prototype_feature_vector)

```

Created prototype feature vector with 10 dimensions (originally 777 dimensions)

Prototype vector tags (sorted by occurrence in filtered_tags): guitar, electric, acoustic, distorted, chord, distortion, Guitar, rock, loop, riff

Created prototype feature vector with 10 dimensions (originally 1006 dimensions)

Prototype vector keywords (sorted by occurrence in filtered_keywords): guitar, electric guitar, acoustic guitar, Recorded, played, Guitar, loop, playing, distorted, distortion

Final feature vector size 15 features: [u'Guitar', u'chord', u'electric', u'acoustic guitar', u'played', u'electric guitar', u'rock', u'Recorded', u'distortion', u'guitar', u'riff', u'distorted', u'acoustic', u'playing', u'loop']

4) Project documents in the vector space

```
In [6]: # Example of getting feature vector from tags list...
random_sound = random.choice(dataset[random.choice(dataset.keys())])
random_sound_tags = random_sound['tags']
random_sound_keywords = random_sound['keywords']
random_sound_features = list(set(random_sound_tags) | set(random_sound_keywords) - set(NO_KEYWORDS))
random_sound_feature_vector = utils.get_feature_vector_from_tags(random_sound_features, prototype_feature_vector)

display(HTML(utils.get_sound_embed_html(random_sound['id'])))
print 'Chosen sound has features:', ', '.join(random_sound_features)
html = '<table><tr><td>'
html += '</td><td>'.join([f for f in prototype_feature_vector])
html += '</td></tr><tr><td>'
html += '</td><td>'.join(['x' if prototype_feature_vector[count] in random_sound_features else ''
                           for count, item in enumerate(prototype_feature_vector)])
html += '</td></tr></table>'
display(HTML(html))
```

nate mess up edit.wav

-00:08



zapama

bend	118bpm	118
distorted	bpm	blues

Chosen sound has features: chopped, sample, 118bpm, electric, Nate, level, BPM, distorted, bpm, 118, guitar, riff, bend, rock, loops, overdrive, blues, loop

Guitar	chord	electric	acoustic guitar	played	electric guitar	rock	Recorded	distortion	guitar	riff	distorted	acous
		x				x			x	x	x	

5) Define train and testing set

```
In [7]: def create_train_and_test_sets(dataset, class_names, n_sounds_per_class,
percentage_training_data,
max_input_tags_for_testing):

    n_training_sounds_per_class = int(n_sounds_per_class * percentage_training_data)
    training_set = dict()
    testing_set = dict()

    # Get 'n_training_sounds_per_class' sounds per class
    for class_name, sounds in dataset.items():
        sounds_from_class = sounds[:] # Copy the list so when we later shuffle it does not affect the original data
        #random.shuffle(sounds_from_class)
        training_set[class_name] = sounds_from_class[:n_training_sounds_per_class] # First sounds for training
        testing_set[class_name] = sounds_from_class[n_training_sounds_per_class:] # Following sounds for testing

        # Save a trimmed version of input tags for testing sounds
        for sound in testing_set[class_name]:
            sound['tags'] = random.sample(sound['tags'], min(max_input_tags_for_testing, len(sound['tags'])))

    print 'Created training and testing sets with the following number of sounds:\n\tTrain\tTest'
    for class_name in class_names:
        training_sounds = training_set[class_name]
        testing_sounds = testing_set[class_name]
        print '\t%i\t%i\t%s' % (len(training_sounds), len(testing_sounds), class_name)
    return training_set, testing_set

training_set, testing_set = create_train_and_test_sets(
    dataset=dataset,
    class_names=CLASS_NAMES,
    n_sounds_per_class=N,
    percentage_training_data=PERCENTAGE_OF_TRAINING_DATA,
    max_input_tags_for_testing=MAX_INPUT_TAGS_FOR_TESTING,
)
```

Created training and testing sets with the following number of sounds:

Train	Test	
50	50	Distorted Guitar
50	50	Electric Guitar
50	50	Acoustic Guitar

6) Train classifier

```
In [8]: def build_tag_feature_vector(sound):
        features = list(set(sound['tags']) | set(sound['keywords']) - set(NO
        _KEYWORDS))
        tag_features = utils.get_feature_vector_from_tags(features, prototyp
        e_feature_vector)
        return np.concatenate([], tag_features))

def train_classifier(training_set, classifier_type, class_names, dataset
_name, feature_vector_func,
                    feature_vector_dimension_labels=None):

    # Prepare data for fitting classifier (as sklearn classifiers requir
    e)
    classes_vector = list()
    feature_vectors = list()
    for class_name, sounds in training_set.items():
        for count, sound in enumerate(sounds):
            # Use index of class name in CLASS_NAMES as numerical value
            (classifier internally represents
            # class label as number)
            classes_vector.append(CLASS_NAMES.index(class_name))
            feature_vector = feature_vector_func(sound)
            if DESCRIPTORS:
                try:
                    mfcc = sound['analysis']['lowlevel']['mfcc']['mean']
                    mfcc_norm = [float(i)/sum(mfcc) for i in mfcc[1:]]
                    feature_vector = np.append(feature_vector, mfcc_norm
                )
            except:
                feature_vector = np.append(feature_vector, [0,0,0,0,
                0,0,0,0,0,0,0,0,0])
            feature_vectors.append(feature_vector)
    # Create and fit classifier
    print 'Training classifier (%s) with %i sounds...' % (CLASSIFIER_TYP
    E, len(feature_vectors)),
    if classifier_type == 'svm':
        classifier = svm.LinearSVC()
        classifier.fit(feature_vectors, classes_vector)
    elif classifier_type == 'tree':
        classifier = tree.DecisionTreeClassifier(max_depth=5)
        classifier.fit(feature_vectors, classes_vector)
        # Plot classifier decision rules
        # WARNING: do not run this if tree is too big, might freeze
        out_filename = os.getcwd() + '/%s-tree.png' % dataset_name
        #utils.export_tree_as_graph(
        #    classifier, feature_vector_dimension_labels, class_names=cl
        ass_names, filename=out_filename)
        #display(HTML('<h4>Learned tree:</h4>' % out_file
        name))
    else:
        raise Exception('Bad classifier type!!!')
    return classifier

classifier = train_classifier(
    training_set=training_set,
    classifier_type=CLASSIFIER_TYPE,
    class_names=CLASS_NAMES,
    dataset_name=DATASET_NAME,
    feature_vector_func=build_tag_feature_vector,
    feature_vector_dimension_labels=prototype_feature_vector,
)
```

Training classifier (tree) with 150 sounds...

7) Evaluate classification

```

In [9]: def evaluate_classifier(testing_set, classifier, class_names, feature_v
        ctor_func, show_confussing_matrix=True):
        # Test with testing set
        print '\nEvaluating with %i instances...' % sum([len(sounds) for sou
nds in testing_set.values()]),
        predicted_data = list()
        for class_name, sounds in testing_set.items():
            for count, sound in enumerate(sounds):
                feature_vector = feature_vector_func(sound)
                if DESCRIPTORS:
                    try:
                        mfcc = sound['analysis']['lowlevel']['barkbands']['m
                        mfcc_norm = [float(i)/sum(mfcc) for i in mfcc[1:]]
                        feature_vector = np.append(feature_vector.append, mf
                        cc_norm)
                    except:
                        feature_vector = np.append(feature_vector, [0,0,0,0,
                        0,0,0,0,0,0,0,0,0])

                        predicted_class_name = unicode(class_names[classifier.predic
t([feature_vector])[0]])
                        predicted_data.append((sound['id'], class_name, predicted_cl
ass_name))
                print 'done!'

        # Compute overall accuracy
        good_predictions = len([1 for sid, cname, pname in predicted_data if
cname == pname])
        wrong_predictions = len([1 for sid, cname, pname in predicted_data i
f cname != pname])
        print '%i correct predictions' % good_predictions
        print '%i wrong predictions' % wrong_predictions
        accuracy = float(good_predictions)/(good_predictions + wrong_predict
ions)
        print 'Overall accuracy %.2f%%' % (100 * accuracy)

        if show_confussing_matrix:
            # Compute confussion matrix (further analysis)
            matrix = list()
            for class_name in CLASS_NAMES:
                predicted_classes = list()
                for sid, cname, pname in predicted_data:
                    if cname == class_name:
                        predicted_classes.append(pname)
                matrix.append([predicted_classes.count(target_class) for tar
get_class in CLASS_NAMES])

            # Plot confussion matrix
            fig = plt.figure()
            plt.clf()
            ax = fig.add_subplot(111)
            ax.set_aspect(1)
            res = ax.imshow(matrix, cmap=plt.cm.Blues, interpolation='neares
t')

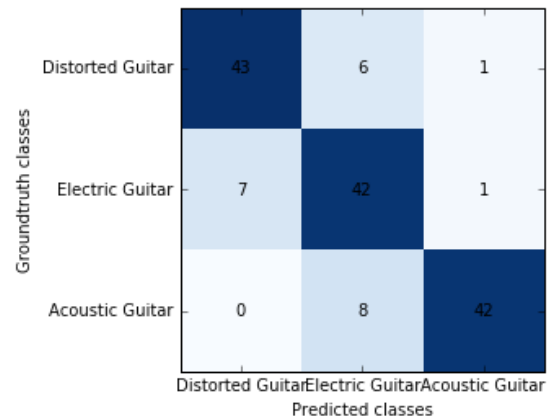
            for x in xrange(len(matrix)):
                for y in xrange(len(matrix)):
                    ax.annotate(str(matrix[x][y]), xy=(y, x),
                                horizontalalignment='center',
                                verticalalignment='center')

            plt.xticks(range(len(class_names)), class_names)
            plt.yticks(range(len(class_names)), class_names)
            plt.xlabel('Predicted classes')
            plt.ylabel('Groundtruth classes')

            print 'Confussion matrix'

```

Evaluating with 150 instances... done!
127 correct predictions
23 wrong predictions
Overall accuracy 84.67%
Confusion matrix



Out[9]: 0.8466666666666667

```
In [10]: # Run the whole experiment again with other parameters

NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR = 10
PERCENTAGE_OF_TRAINING_DATA = 0.75
MAX_INPUT_TAGS_FOR_TESTING = 2

for classifier_type in ['svm', 'tree']:
    display(HTML('<h2>With %s classifier</h2>' % classifier_type.upper()))

    prototype_feature_vector = build_tag_vector_space(
        n_dimensions=NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR,
        dataset=dataset,
        class_names=CLASS_NAMES,
    )
    training_set, testing_set = create_train_and_test_sets(
        dataset=dataset,
        class_names=CLASS_NAMES,
        n_sounds_per_class=N,
        percentage_training_data=PERCENTAGE_OF_TRAINING_DATA,
        max_input_tags_for_testing=MAX_INPUT_TAGS_FOR_TESTING,
    )
    classifier = train_classifier(
        training_set=training_set,
        classifier_type=classifier_type,
        class_names=CLASS_NAMES,
        dataset_name=DATASET_NAME,
        feature_vector_func=build_tag_feature_vector,
        feature_vector_dimension_labels=prototype_feature_vector,
    )
    evaluate_classifier(
        testing_set=testing_set,
        classifier=classifier,
        class_names=CLASS_NAMES,
        feature_vector_func=build_tag_feature_vector,
```


With SVM classifier

Created prototype feature vector with 10 dimensions (originally 777 dimensions)

Prototype vector tags (sorted by occurrence in filtered_tags): guitar, electric, acoustic, distorted, chord, distortion, Guitar, rock, loop, riff

Created training and testing sets with the following number of sounds:

Train	Test	
75	25	Distorted Guitar
75	25	Electric Guitar
75	25	Acoustic Guitar

Training classifier (tree) with 225 sounds...

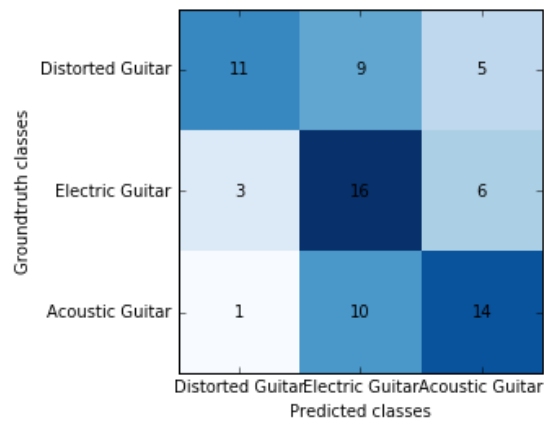
Evaluating with 75 instances... done!

41 correct predictions

34 wrong predictions

Overall accuracy 54.67%

Confussion matrix



With TREE classifier

Created prototype feature vector with 10 dimensions (originally 683 dimensions)

Prototype vector tags (sorted by occurrence in filtered_tags): guitar, electric, acoustic, distorted, chord, Guitar, rock, distortion, loop, riff

Created training and testing sets with the following number of sounds:

Train	Test	
75	25	Distorted Guitar
75	25	Electric Guitar
75	25	Acoustic Guitar

Training classifier (tree) with 225 sounds...

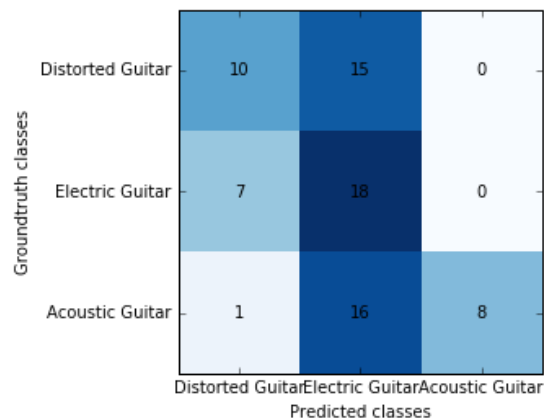
Evaluating with 75 instances... done!

36 correct predictions

39 wrong predictions

Overall accuracy 48.00%

Confussion matrix



```
In [11]: # CODE FOR GENERATING FEATURE VECTORS WITH LATENT SEMANTIC ANALYSIS

def build_latent_semantic_analysis_feature_vectors(n_dimensions, training_set, testing_set):
    # Build latent semantic analysis (LSA) matrix factors
    vectorizer = TfidfVectorizer(
        sublinear_tf=True, max_df=0.5, stop_words='english', analyzer='word', max_features=50000)
    svd = TruncatedSVD(n_components=n_dimensions, n_iter=7, random_state=42)

    # Compute LSA feature vectors for training set
    lsa_feature_vectors = list()
    lsa_feature_vectors_sound_ids = list()
    for class_name, sounds in training_set.items():
        for sound in sounds:
            lsa_feature_vectors.append(' '.join(sound['tags']))
            lsa_feature_vectors_sound_ids.append(sound['id'])
    X = vectorizer.fit_transform(lsa_feature_vectors)
    lsa_feature_vectors = svd.fit_transform(X)
    lsa_feature_vectors_dict = {sound_id:features for sound_id, features
                               in zip(lsa_feature_vectors_sound_ids, lsa_feature_vectors)}

    # Now compute LSA feature vectors for testing set (using vectorizer and svd learnt for training set)
    lsa_feature_vectors = list()
    lsa_feature_vectors_sound_ids = list()
    for class_name, sounds in testing_set.items():
        for sound in sounds:
            lsa_feature_vectors.append(' '.join(sound['tags']))
            lsa_feature_vectors_sound_ids.append(sound['id'])
    X = vectorizer.transform(lsa_feature_vectors)
    lsa_feature_vectors = svd.transform(X)
    lsa_feature_vectors_dict.update({sound_id:features for sound_id, features
                                     in zip(lsa_feature_vectors_sound_ids, lsa_feature_vectors)})

    print 'Built LSA feature vectors for %i sounds!' % len(lsa_feature_vectors_dict)
    return lsa_feature_vectors_dict
```

```

In [12]: # EXAMPLE EXPERIMENT USING LATENT SEMANTIC ANALYSIS

display(HTML('<h2>With Latent Semantic Analysis</h2>'))

NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR = 10
CLASSIFIER_TYPE = 'tree'
PERCENTAGE_OF_TRAINING_DATA = 0.75
MAX_INPUT_TAGS_FOR_TESTING = 2

def build_lsa_feature_vector(sound):
    lsa_features = lsa_feature_vectors_dict[sound['id']]
    return np.concatenate([[], lsa_features])

lsa_feature_vectors_dict = build_latent_semantic_analysis_feature_vectors(
    n_dimensions=NUMBER_OF_DIMENSIONS_OF_FEATURE_VECTOR, training_set=train_
    aining_set, testing_set=testing_set,)

training_set, testing_set = create_train_and_test_sets(
    dataset=dataset, class_names=CLASS_NAMES, n_sounds_per_class=N,
    percentage_training_data=PERCENTAGE_OF_TRAINING_DATA, max_input_tags
    _for_testing=MAX_INPUT_TAGS_FOR_TESTING)
classifier = train_classifier(
    training_set=training_set, classifier_type=CLASSIFIER_TYPE, class_na
    mes=CLASS_NAMES, dataset_name=DATASET_NAME,
    feature_vector_func=build_lsa_feature_vector,
)
evaluate_classifier(
    testing_set=testing_set, classifier=classifier, class_names=CLASS_NA
    MES,
    feature_vector_func=build_lsa_feature_vector)

```

With Latent Semantic Analysis

Built LSA feature vectors for 291 sounds!
 Created training and testing sets with the following number of sounds:

Train	Test	
75	25	Distorted Guitar
75	25	Electric Guitar
75	25	Acoustic Guitar

Training classifier (tree) with 225 sounds...

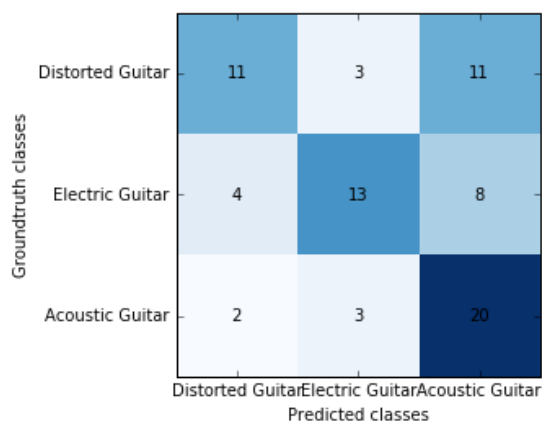
Evaluating with 75 instances... done!

44 correct predictions

31 wrong predictions

Overall accuracy 58.67%

Confussion matrix



Out[12]: 0.5866666666666667

In []: